

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("uber.csv")

df.head()
df.info() #To get the required information of the dataset
df.columns #TO get number of columns in the dataset
df = df.drop(['Unnamed: 0', 'key'], axis= 1) #To drop unnamed c
df.head()
df.shape #To get the total (Rows,Columns)
df.dtypes #To get the type of each column
df.info()
df.describe() #To get statistics of each columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8918 entries, 0 to 8917
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            8918 non-null   int64
1   key                   8918 non-null   object
2   fare_amount           8917 non-null   float64
3   pickup_datetime      8917 non-null   object
4   pickup_longitude      8917 non-null   float64
5   pickup_latitude       8917 non-null   float64
6   dropoff_longitude     8917 non-null   float64
7   dropoff_latitude      8917 non-null   float64
8   passenger_count       8917 non-null   float64
dtypes: float64(6), int64(1), object(2)
memory usage: 627.2+ KB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8918 entries, 0 to 8917
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fare_amount           8917 non-null   float64
1   pickup_datetime      8917 non-null   object
2   pickup_longitude      8917 non-null   float64
3   pickup_latitude       8917 non-null   float64
4   dropoff_longitude     8917 non-null   float64
5   dropoff_latitude      8917 non-null   float64
6   passenger_count       8917 non-null   float64
```

```

0    passenger_count    8917 non-null    float64
dtypes: float64(6), object(1)
memory usage: 487.8+ KB

```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_la
count	8917.00000	8917.000000	8917.000000	8917.000000	8917.
mean	11.44521	-72.700539	40.000828	-72.673399	40.
std	10.41732	12.320446	5.895643	9.855862	5.
min	2.50000	-748.016667	-74.009697	-75.350437	-73.
25%	6.00000	-73.992015	40.734997	-73.991472	40.
50%	8.50000	-73.981582	40.752407	-73.979908	40.
75%	12.50000	-73.967155	40.767058	-73.963588	40.
max	350.00000	40.770667	41.366138	40.761672	41.

```
df.isnull().sum()
```

```

fare_amount      1
pickup_datetime  1
pickup_longitude  1
pickup_latitude  1
dropoff_longitude 1
dropoff_latitude  1

```

```
passenger_count    1  
dtype: int64
```

[+ Code](#)[+ Text](#)

```
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean  
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].me
```

```
df.isnull().sum()
```

```
fare_amount        1  
pickup_datetime   1  
pickup_longitude   1  
pickup_latitude    1  
dropoff_longitude   0  
dropoff_latitude    0  
passenger_count    1  
dtype: int64
```

```
df.dtypes
```

```
fare_amount        float64  
pickup_datetime     object  
pickup_longitude    float64
```

```
pickup_latitude      float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       float64
dtype: object
```

```
df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors=
df.dtypes
```

```
fare_amount           float64
pickup_datetime       datetime64[ns, UTC]
pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       float64
dtype: object
```

```
df= df.assign(hour = df.pickup_datetime.dt.hour,
day= df.pickup_datetime.dt.day,
month = df.pickup_datetime.dt.month,
year = df.pickup_datetime.dt.year,
```

```
dayofweek = df.pickup_datetime.dt.dayofweek)
df.head()
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude
0	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.738354	-73.99951
1	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.728225	-73.99471
2	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.740770	-73.96256

```
df = df.drop('pickup_datetime',axis=1)
```

```
df.dtypes
```

```
fare_amount          float64
pickup_longitude      float64
pickup_latitude       float64
dropoff_longitude     float64
dropoff_latitude      float64
passenger_count       float64
```


```
hour          float64
day           float64
month         float64
year          float64
dayofweek     float64
dtype: object
```

```
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15
```


```

fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count  AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour             AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day              AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month            AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year             AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek        AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object


```




fare_amount




pickup_longitude

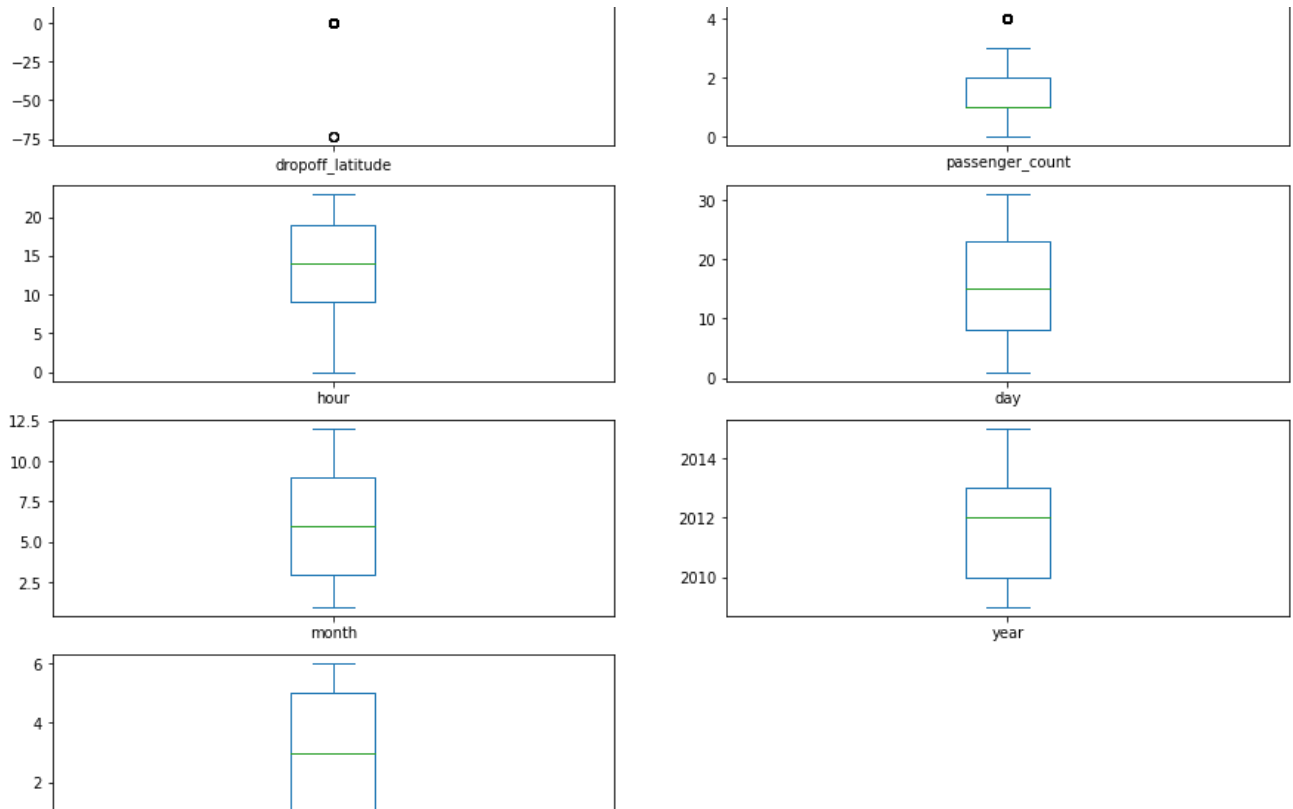


pickup_latitude



dropoff_longitude





```
def remove_outlier(df1 , col):  
    Q1 = df1[col].quantile(0.25)
```

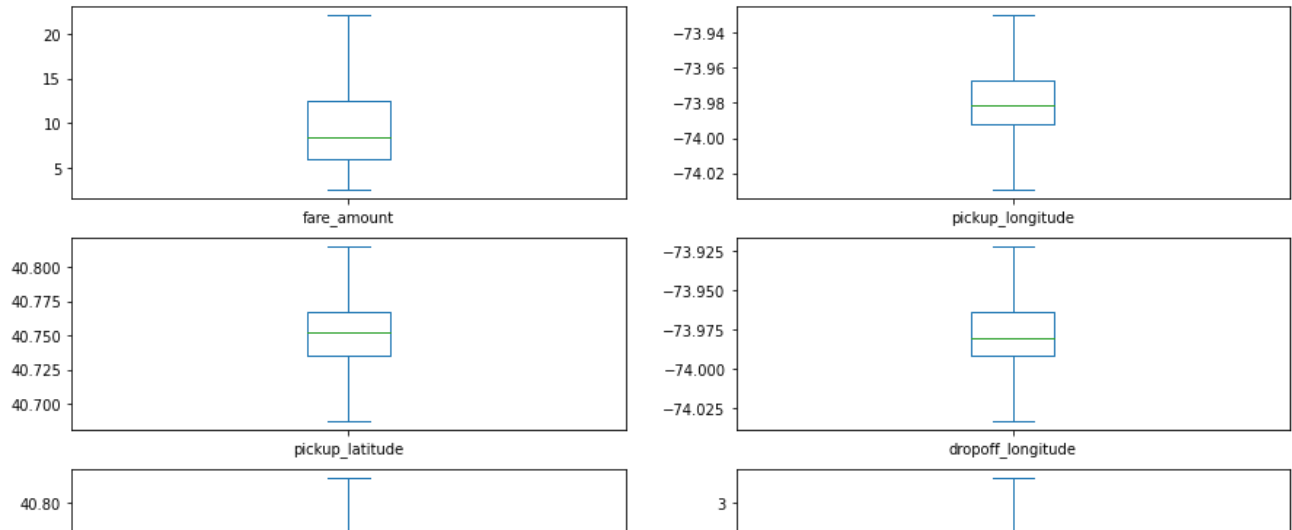
```
Q3 = df1[col].quantile(0.75)
IQR = Q3 - Q1
lower_whisker = Q1-1.5*IQR
upper_whisker = Q3+1.5*IQR
df[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
return df1








def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(df , c)
    return df1
df = treat_outliers_all(df , df.iloc[:, 0::])
df.plot(kind = "box",subplots = True,layout = (7,2),figsize=(15
```

```

fare_amount      AxesSubplot(0.125,0.787927;0.352273x0.0920732)
pickup_longitude AxesSubplot(0.547727,0.787927;0.352273x0.0920732)
pickup_latitude  AxesSubplot(0.125,0.677439;0.352273x0.0920732)
dropoff_longitude AxesSubplot(0.547727,0.677439;0.352273x0.0920732)
dropoff_latitude AxesSubplot(0.125,0.566951;0.352273x0.0920732)
passenger_count  AxesSubplot(0.547727,0.566951;0.352273x0.0920732)
hour             AxesSubplot(0.125,0.456463;0.352273x0.0920732)
day              AxesSubplot(0.547727,0.456463;0.352273x0.0920732)
month            AxesSubplot(0.125,0.345976;0.352273x0.0920732)
year             AxesSubplot(0.547727,0.345976;0.352273x0.0920732)
dayofweek        AxesSubplot(0.125,0.235488;0.352273x0.0920732)
dtype: object


```





```
pip install haversine
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Collecting haversine
  Downloading haversine-2.7.0-py2.py3-none-any.whl (6.9 kB)
Installing collected packages: haversine
Successfully installed haversine-2.7.0
```



```
import haversine as hs
```

```
travel_dist = []
for pos in range(len(df['pickup_longitude'])):
    long1,lati1,long2,lati2 = [df['pickup_longitude'][pos],df['
    loc1=(lati1,long1)
    loc2=(lati2,long2)
```


```

c = hs.haversine(loc1,loc2)
travel_dist.append(c)
print(travel_dist)
df['dist_travel_km'] = travel_dist
df.head()

```

```
[1.6833250775073447, 2.4575932783467835, 5.036384146783453, 1.661685753650294, 4.1102]
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	-73.999817	40.738354	-73.999512	40.7232
1	7.7	-73.994355	40.728225	-73.994710	40.7503
2	12.9	-74.005043	40.740770	-73.962565	40.7726
3	5.3	-73.976124	40.790844	-73.965316	40.8033
4	16.0	-73.929865	40.744085	-73.973082	40.7612



```

df= df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)
print("Remaining observastions in the dataset:", df.shape)

```

Remaining observations in the dataset: (8917, 12)

```
incorrect_coordinates = df.loc[(df.pickup_latitude>90)|(df.pick  
(df.dropoff_latitude>90)|(df.dropoff_latitude<-90)|  
(df.pickup_longitude>180)|(df.pickup_longitude<-180)|  
(df.dropoff_longitude>90)|(df.dropoff_longitude<-90)  
]
```

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore'  
df.head()
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:4913: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/usage>

```
df.drop(incorrect_coordinates, inplace = True, errors = 'ignore')
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:4913: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/usage>

```
errors=errors,
```



```
df.head()
```


fare_amount pickup_longitude pickup_latitude dropoff_longitude dropoff_latitu

```
df.isnull().sum()
```

```
fare_amount      0
pickup_longitude  0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
passenger_count   0
hour              0
day               0
month             0
year              0
dayofweek         0
dist_travel_km    0
dtype: int64
```

```
sns.heatmap(df.isnull())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8ab5cbc6d0>
```



```
corr=df.corr() #function to find correlation
```

```
ck  pk  cp  fr  ss
```

```
u
```

```
fig,axis = pl.subplots(figsize = (10,6))
```



```
sns.heatmap(df.corr(),annot=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8ab3216b10>
```



```
x = df[['pickup_longitude','pickup_latitude','dropoff_longitude',
y = df['fare_amount']
```

```
    x = df[['pickup_longitude','pickup_latitude','dropoff_longitude',
    y = df['fare_amount']
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size
```

```
from sklearn.linear_model import LinearRegression
```

```
regression = LinearRegression()
regression.fit(X_train,y_train)
regression.coef_ #To find the linear coeeficient
regression.intercept_ #To find the linear intercept
prediction = regression.predict(X_test) #To predict the target
print(prediction)
y_test
```

```
[ 8.15350341 11.45372045 20.88795037 ...  4.57442819  8.94565272
 6.08170728]
3731      8.10
8472     22.25
8255     21.70
8776      7.30
5638      5.30
...
7491      5.50
3294      6.90
1833      3.30
6797      9.30
481       5.30
Name: fare_amount, Length: 2943, dtype: float64
```

```
regression = LinearRegression()
regression.fit(X_train,y_train)
regression.coef_ #To find the linear coeeficient
regression.intercept_ #To find the linear intercept
prediction = regression.predict(X_test) #To predict the target
print(prediction)
y_test
```

```
[ 8.15350341 11.45372045 20.88795037 ...  4.57442819  8.94565272
 6.08170728]
3731      8.10
8472     22.25
8255     21.70
8776      7.30
5638      5.30
...
7491      5.50
3294      6.90
1833      3.30
6797      9.30
481       5.30
Name: fare_amount, Length: 2943, dtype: float64
```

```
from sklearn.metrics import r2_score

r2_score(y_test,prediction)
from sklearn.metrics import mean_squared_error
MSE = mean_squared_error(y_test,prediction)
MSE
RMSE = np.sqrt(MSE)
RMSE
```

```
3.0423611997128597
```

```
from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=100)
rf.fit(X_train,y_train)
y_pred = rf.predict(X_test)
y_pred
```

```
array([ 8.099, 13.686, 21.704, ...,  4.451,  8.822,  4.039])
```

```
R2_Random = r2_score(y_test,y_pred)
```

```
R2_Random
```

```
0.7917656690094285
```

```
MSE_Random = mean_squared_error(y_test,y_pred)
```

```
MSE_Random
```

```
6.286907173010533
```

```
RMSE_Random = np.sqrt(MSE_Random)
```

```
RMSE_Random
```

```
2.507370569543029
```


[Colab paid products](#) - [Cancel contracts here](#)

